

State Space Models for Time Series

Architecture, Mathematics, and Implementation

Senior ML Engineer

Wageningen University & Research



January 21, 2026

Overview

- 1 Motivation
- 2 Encoder Architecture
- 3 The Mamba Block
- 4 SSM Mathematics
- 5 Data Processing (Recurrence Plots)
- 6 Algorithm

Why State Space Models? (1/3)

Challenges in Time Series Modeling



Why State Space Models? (2/3)

State Space Model Solutions

Solution: Continuous State Memory

SSM maintains hidden state h_t

$$h_t = f(A, B, h_{t-1}, x_t)$$

Benefits:

- Compresses full history efficiently
- Constant memory complexity
- Captures long-range patterns

Solution: Efficient Sequential Scan

Sequential processing enables:

Linear time: $O(T)$
Parallel training
Hardware-optimized

Result:


*Scales to very long sequences!
(thousands of time steps)*

Why State Space Models? (3/3)

Key Advantages of SSMs:

- **Efficient Memory**: Compresses long history into fixed-size state
 - Maintains hidden state h_t that encodes full sequence history
 - Constant memory complexity regardless of sequence length
- **Linear Complexity**: $O(T)$ vs $O(T^2)$ for attention
 - Sequential processing enables efficient scanning
 - Scales to very long sequences (thousands of time steps)
- **Long Dependencies**: Designed for sequences with distant correlations
 - HiPPO initialization optimizes for long-range memory
 - Captures patterns across entire time series

Our Running Example (1/2)




Our Running Example (2/2)

Throughout this presentation:


- We'll follow a **univariate time series** of length $T = 96$
 - Real-world time series with trend, seasonality, and noise
 - Representative of typical forecasting scenarios
- Focus on the **first 2 values**: $x_0 = 0.5$ and $x_1 = 0.8$
 - These specific values will appear in all numerical examples
 - Allows tracking data flow through the entire pipeline
- **Goal:** Understand how SSMs process sequential data
 - From raw input to final embedding
 - Step-by-step mathematical transformations

High-Level Encoder Architecture

- **Objective:** Map time series $X \in \mathbb{R}^{B \times T \times F}$ to compact embeddings $E \in \mathbb{R}^{B \times D_{emb}}$.
- **Core Components** (from `mamba_encoder.yaml`):
 - ① **Tokenization:** Slicing time series into windows.
 - ② **Projection:** Mapping tokens to model dimension D_{model} .
 - ③ **Backbone:** Stack of L Mamba Blocks (SSM).
 - ④ **Pooling:** Aggregating sequence info (Mean, Last, or CLS).



Encoder Variants (1/2)



Encoder Variants (2/2)

Standard Encoder

- **Input:** Raw sequence values
- **Patching:** Direct 1D windows
- **Projection:** Linear layer (`nn.Linear`)
- **Processing:** Sequential SSM blocks
- **Use Case:** General time series forecasting

Visual Encoder

- **Input:** Time series → pseudo-images
- **Transform:** Recurrence Plot
 - $R_{ij} = \|x_i - x_j\|$
- **Projection:** 2D convolution (`Conv2d`)
- **Processing:** Vision-style SSM
- **Use Case:** Capturing structural patterns

Inside the Mamba Block

The MambaBlock handles the sequence mixing.

Forward Pass (Simplified):

- ① **Input Projection:** $x \in \mathbb{R}^{B \times T \times D} \rightarrow z, x' \in \mathbb{R}^{B \times T \times E}$
- ② **Convolution:** 1D causal conv on x'
- ③ **SSM Processing:** `_selective_scan(x' , Δ , A , B , C)`
- ④ **Output Projection:** Combine with gating z and project back


Selective Parameters: Δ, B, C are *input-dependent*. This is the core innovation!

Step 1: Input Sequence

The time series x enters the model. Each time step is a scalar (univariate) or vector (multivariate).

$$\mathbf{u} : \mathbb{R} \rightarrow \mathbb{R}^F$$

For our example: $u_0 = 0.5, u_1 = 0.8, \dots$




Step 2: Latent State $h(t)$

SSM maintains a **continuous hidden state** that captures history.

$$\frac{dh}{dt} = \mathbf{A}h(t) + \mathbf{B}u(t)$$

A encodes state dynamics (HiPPO-initialized).
B maps input influence.



Discretization (Zero-Order Hold)

To process sampled data with dynamic step sizes Δ , we discretize the continuous system. Given a step size Δ , the discrete parameters $\bar{\mathbf{A}}$ and $\bar{\mathbf{B}}$ are:

$$\bar{\mathbf{A}} = \exp(\Delta \mathbf{A}) \quad (1)$$

$$\begin{aligned} \bar{\mathbf{B}} &= (\Delta \mathbf{A})^{-1} (\exp(\Delta \mathbf{A}) - \mathbf{I}) \cdot \Delta \mathbf{B} \\ &= \mathbf{A}^{-1} (\bar{\mathbf{A}} - \mathbf{I}) \cdot \mathbf{B} \end{aligned} \quad (2)$$

This matches the specific implementation in `mamba_block.py`:

```
integral = torch.linalg.solve(A_expand, A_expm - eye)
B_disc = torch.bmm(integral, B_expand)
```

Section Summary: SSM Mathematics (1/2)

Continuous System

$$\frac{dh}{dt} = \mathbf{A}h + \mathbf{B}u \quad (\text{Differential equation})$$

Models smooth evolution of hidden state over continuous time

Discretization (Zero-Order Hold)

$$\bar{\mathbf{A}} = \exp(\Delta \mathbf{A}), \quad \bar{\mathbf{B}} = \mathbf{A}^{-1}(\bar{\mathbf{A}} - \mathbf{I})\mathbf{B}$$

Converts continuous dynamics to discrete time steps of size Δ


Discrete Recurrence

$$h_t = \bar{\mathbf{A}}h_{t-1} + \bar{\mathbf{B}}u_t \quad (\text{State update})$$

$$y_t = \mathbf{C}h_t \quad (\text{Output projection})$$

Enables efficient sequential processing of sampled data

Section Summary: SSM Mathematics (2/2)




Key Concept: Zero-Order Hold (ZOH) maintains input constant between samples

Step 3: Time Series Patching

The raw time series is typically long and noisy.

We extract **overlapping patches** of fixed length (e.g., 96 steps).

This is analogous to image patching in Vision Transformers.




Step 4: Feature Extraction

From each patch, we can extract:

- **Raw values:** Direct encoding
- **Statistical features:** Mean, variance, etc.
- **Structural features:** Recurrence structure

The Visual Encoder uses the last option.



Step 5: Visual Transformation (RP)


We convert the 1D patch into a 2D Recurrence Plot.

$$R_{i,j} = \|\mathbf{x}_i - \mathbf{x}_j\|$$

Numerical Example: Sequence patch (first 2 values): $\mathbf{x} = [0.5, 0.8]$

$$\mathbf{R} = \begin{pmatrix} |0.5 - 0.5| & |0.5 - 0.8| \\ |0.8 - 0.5| & |0.8 - 0.8| \end{pmatrix}$$

$$\mathbf{R} = \begin{pmatrix} 0 & 0.3 \\ 0.3 & 0 \end{pmatrix}$$



Algorithm: Selective Scan

Logic verified in: `src/models/mamba_block.py`


Algorithm 1 Selective Scan with HiPPO-Initialized SSM

Require: Input $\mathbf{x} \in \mathbb{R}^{B \times T \times D}$, step sizes $\boldsymbol{\delta} \in \mathbb{R}^{B \times T \times 1}$

Ensure: Output $\mathbf{y} \in \mathbb{R}^{B \times T \times D}$

- 1: Precompute $\mathbf{A}, \mathbf{B}, \mathbf{C}$ (HiPPO)
 - 2: **for** $t = 0$ **to** $T - 1$ **do**
 - 3: $\mathbf{u}_t \leftarrow \mathbf{x}[:, t, :]$; $\Delta_t \leftarrow \boldsymbol{\delta}[:, t, :]$
 - 4: $\tilde{\mathbf{A}} \leftarrow \Delta_t \cdot \mathbf{A}$ {Scale A by time step}
 - 5: $\mathbf{A}_d \leftarrow \exp(\tilde{\mathbf{A}})$ {Discretize A}
 - 6: $\mathbf{h} \leftarrow \mathbf{A}_d \mathbf{h} + \mathbf{B}_d \mathbf{u}_t$ {Update State}
 - 7: $\mathbf{y}_t \leftarrow \mathbf{C} \mathbf{h}$ {Project to Output}
 - 8: **end for**
-

Numerical Walkthrough: Selective Scan (1/2)



Numerical Walkthrough: Selective Scan (2/2)

Key Insight: At $t = 1$, state h_1 retains memory from both u_0 and u_1 !

- **Input Sequence:**

- $u_0 = 0.5, u_1 = 0.8$ (from our running example)
- Step size: $\Delta = 0.5$

- **State Evolution:**

- $h_0 = \bar{\mathbf{B}}u_0 \approx [0.195, 0.145]^T$
- $h_1 = \bar{\mathbf{A}}h_0 + \bar{\mathbf{B}}u_1 \approx [0.431, 0.318]^T$
- State accumulates information: h_1 depends on both u_0 and u_1

- **Output Sequence:**

- $y_0 = \mathbf{C}h_0 = 0.195$
- $y_1 = \mathbf{C}h_1 = 0.431$

Implementation Specifics

Based on `mamba_block.py`:

- **HiPPO Initialization:** Matrix \mathbf{A} is initialized using the Legendre-S (LegS) measure to handle long-term dependencies.
- **Parallelism vs. Scanning:**
 - The Python implementation performs a **sequential loop** (Line 144 in `mamba_block.py`).
 - Optimized CUDA implementations (not currently used) usually perform a parallel associative scan.
- **Numerical Stability:**
 - Δ is clamped: $\Delta \in [10^{-4}, 3.0]$.
 - `torch.linalg.solve` used instead of explicit inverse for \mathbf{A}^{-1} to ensure stability.

Summary: State Space Models for Time Series

Core Concepts:

- ① **Motivation:** Efficient long-range dependencies
- ② **SSM Math:** Continuous → Discrete via ZOH
- ③ **Selective Scan:** Dynamic state updates
- ④ **Visual Encoding:** Recurrence Plots for 2D patterns

Implementation:

- HiPPO initialization for \mathbf{A}
- Parallel training + Sequential inference
- Two encoder variants: Standard & Visual
- Numerically stable via `torch.linalg.solve`

Running Example: Time series ($T = 96$) with $x_0 = 0.5, x_1 = 0.8$
⇒ State memory: h_1 encodes both past and present!

Questions?