

# State Space Models for Time Series

## Architecture, Mathematics, and Implementation

Leandro Stival

Wageningen University & Research

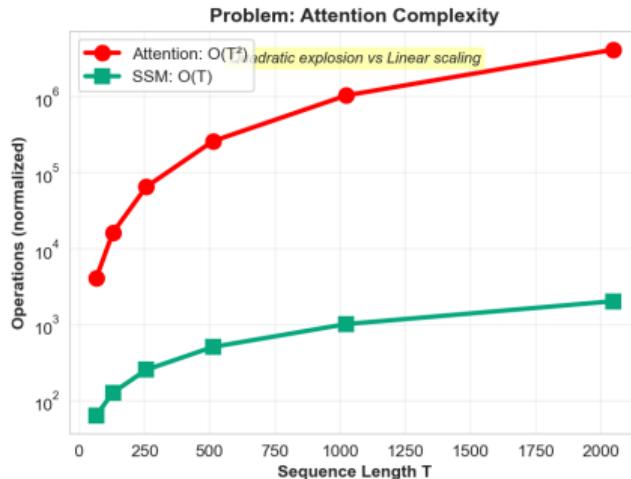
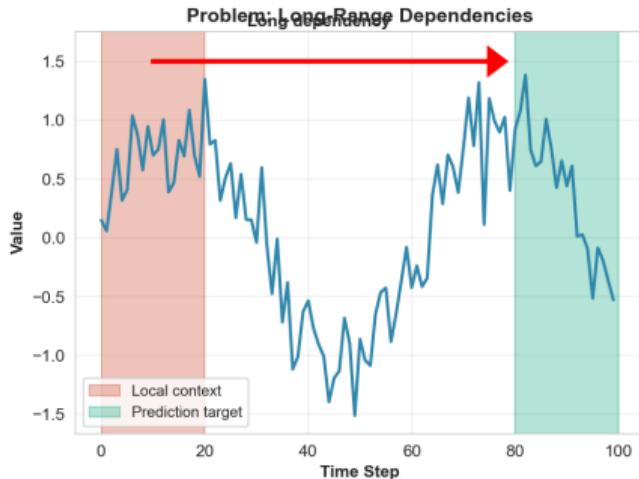
January 21, 2026

# Overview

- 1 Motivation
- 2 Encoder Architecture
- 3 The Mamba Block
- 4 SSM Mathematics
- 5 Data Processing (Recurrence Plots)
- 6 Algorithm

# Why State Space Models? (1/3)

## Challenges in Time Series Modeling



# Why State Space Models? (2/3)

## State Space Model Solutions

Solution: Continuous State Memory

SSM maintains hidden state  $h_t$

$$h_t = f(A, B, h_{t-1}, x_t)$$

Benefits:

- Compresses full history efficiently
- Constant memory complexity
- Captures long-range patterns

Solution: Efficient Sequential Scan

Sequential processing enables:

Linear time:  $O(T)$   
Parallel training  
Hardware-optimized

Result:

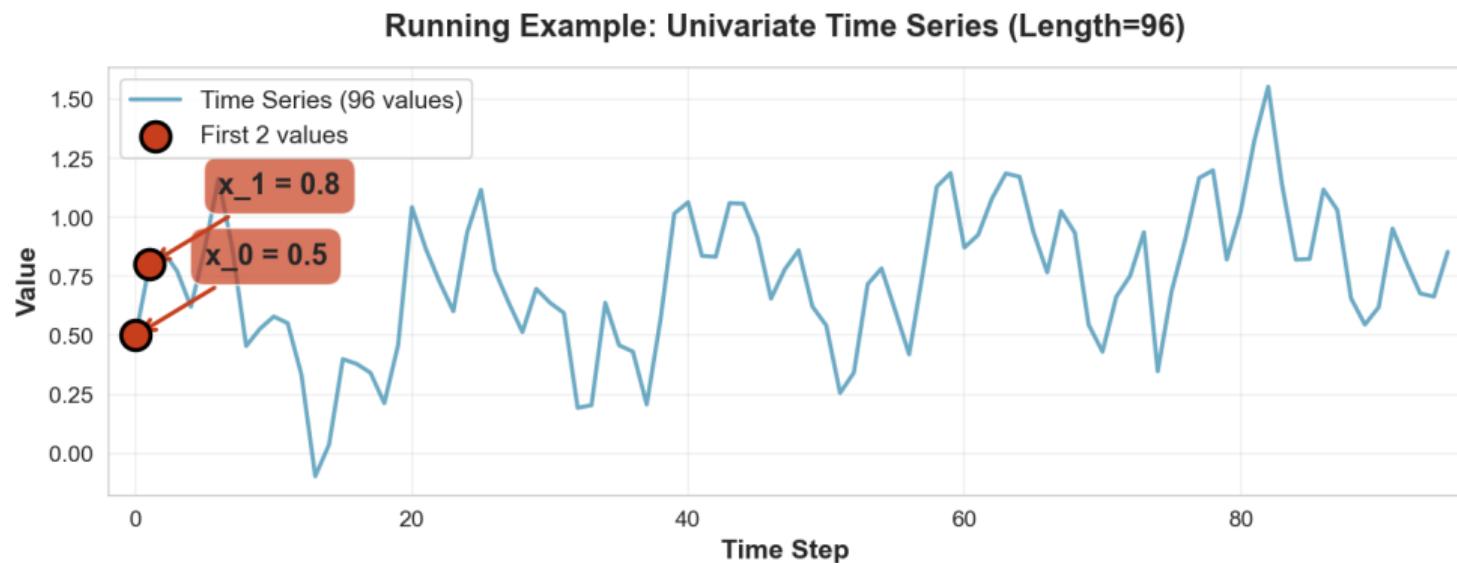
*Scales to very long sequences!  
(thousands of time steps)*

# Why State Space Models? (3/3)

## Key Advantages of SSMs:

- **Efficient Memory**: Compresses long history into fixed-size state
  - Maintains hidden state  $h_t$  that encodes full sequence history
  - Constant memory complexity regardless of sequence length
- **Linear Complexity**:  $O(T)$  vs  $O(T^2)$  for attention
  - Sequential processing enables efficient scanning
  - Scales to very long sequences (thousands of time steps)
- **Long Dependencies**: Designed for sequences with distant correlations
  - HiPPO initialization optimizes for long-range memory
  - Captures patterns across entire time series

# Our Running Example (1/2)



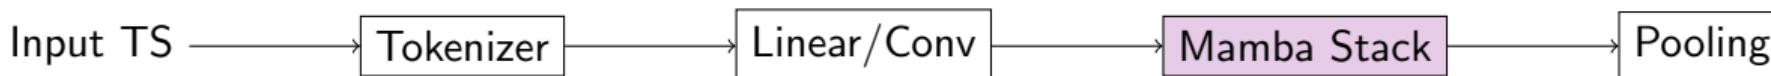
# Our Running Example (2/2)

Throughout this presentation:

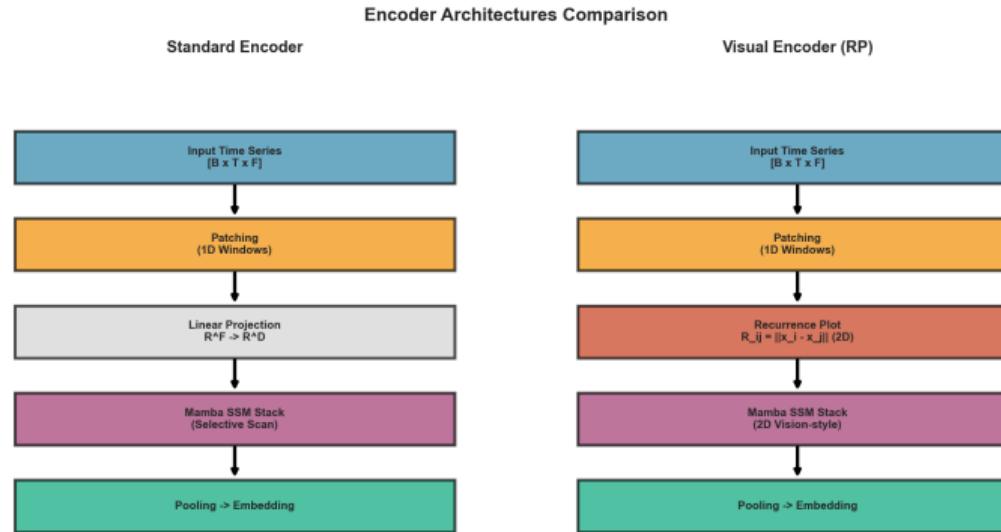
- We'll follow a **univariate time series** of length  $T = 96$ 
  - Real-world time series with trend, seasonality, and noise
  - Representative of typical forecasting scenarios
- Focus on the **first 2 values**:  $x_0 = 0.5$  and  $x_1 = 0.8$ 
  - These specific values will appear in all numerical examples
  - Allows tracking data flow through the entire pipeline
- **Goal:** Understand how SSMs process sequential data
  - From raw input to final embedding
  - Step-by-step mathematical transformations

# High-Level Encoder Architecture

- **Objective:** Map time series  $X \in \mathbb{R}^{B \times T \times F}$  to compact embeddings  $E \in \mathbb{R}^{B \times D_{emb}}$ .
- **Core Components** (from `mamba_encoder.yaml`):
  - ① **Tokenization:** Slicing time series into windows.
  - ② **Projection:** Mapping tokens to model dimension  $D_{model}$ .
  - ③ **Backbone:** Stack of  $L$  Mamba Blocks (SSM).
  - ④ **Pooling:** Aggregating sequence info (Mean, Last, or CLS).



# Encoder Variants (1/2)



# Encoder Variants (2/2)

## Standard Encoder

- **Input:** Raw sequence values
- **Patching:** Direct 1D windows
- **Projection:** Linear layer (`nn.Linear`)
- **Processing:** Sequential SSM blocks
- **Use Case:** General time series forecasting

## Visual Encoder

- **Input:** Time series → pseudo-images
- **Transform:** Recurrence Plot
  - $R_{ij} = \|x_i - x_j\|$
- **Projection:** 2D convolution (`Conv2d`)
- **Processing:** Vision-style SSM
- **Use Case:** Capturing structural patterns

# Inside the Mamba Block

The MambaBlock handles the sequence mixing.

**Forward Pass** (Simplified):

- ① **Input Projection:**  $x \in \mathbb{R}^{B \times T \times D} \rightarrow z, x' \in \mathbb{R}^{B \times T \times E}$
- ② **Convolution:** 1D causal conv on  $x'$
- ③ **SSM Processing:** `_selective_scan( $x'$ ,  $\Delta$ ,  $A$ ,  $B$ ,  $C$ )`
- ④ **Output Projection:** Combine with gating  $z$  and project back

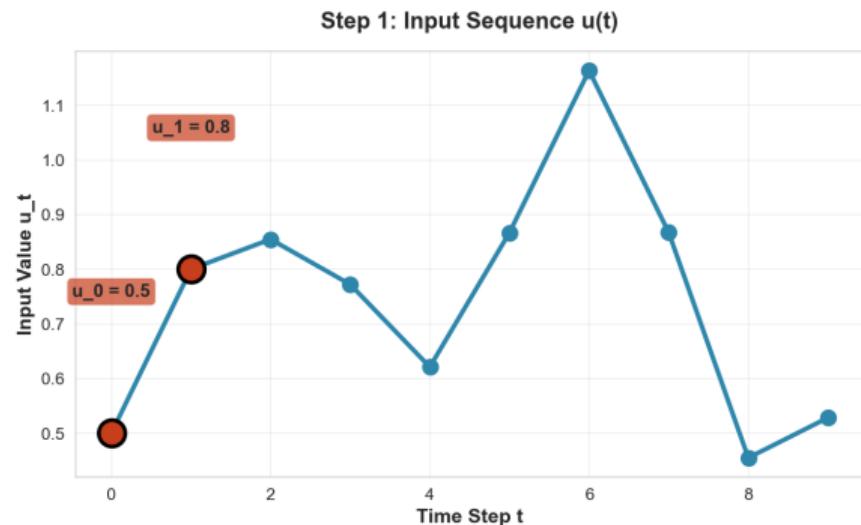
**Selective Parameters:**  $\Delta, B, C$  are *input-dependent*. This is the core innovation!

## Step 1: Input Sequence

The time series  $x$  enters the model. Each time step is a scalar (univariate) or vector (multivariate).

$$\mathbf{u} : \mathbb{R} \rightarrow \mathbb{R}^F$$

For our example:  $u_0 = 0.5, u_1 = 0.8, \dots$

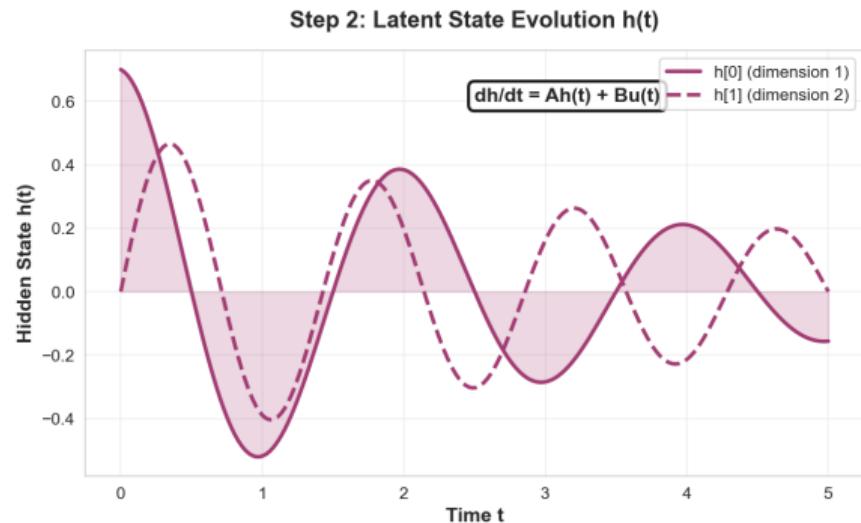


## Step 2: Latent State $h(t)$

SSM maintains a **continuous hidden state** that captures history.

$$\frac{dh}{dt} = \mathbf{A}h(t) + \mathbf{B}u(t)$$

**A** encodes state dynamics (HiPPO-initialized).  
**B** maps input influence.



## Discretization (Zero-Order Hold)

To process sampled data with dynamic step sizes  $\Delta$ , we discretize the continuous system. Given a step size  $\Delta$ , the discrete parameters  $\bar{\mathbf{A}}$  and  $\bar{\mathbf{B}}$  are:

$$\bar{\mathbf{A}} = \exp(\Delta \mathbf{A}) \quad (1)$$

$$\begin{aligned} \bar{\mathbf{B}} &= (\Delta \mathbf{A})^{-1} (\exp(\Delta \mathbf{A}) - \mathbf{I}) \cdot \Delta \mathbf{B} \\ &= \mathbf{A}^{-1} (\bar{\mathbf{A}} - \mathbf{I}) \cdot \mathbf{B} \end{aligned} \quad (2)$$

This matches the specific implementation in `mamba_block.py`:

```
integral = torch.linalg.solve(A_expand, A_expm - eye)
B_disc = torch.bmm(integral, B_expand)
```

## Section Summary: SSM Mathematics (1/2)

### Continuous System

$$\frac{dh}{dt} = \mathbf{A}h + \mathbf{B}u \quad (\text{Differential equation})$$

Models smooth evolution of hidden state over continuous time

### Discretization (Zero-Order Hold)

$$\bar{\mathbf{A}} = \exp(\Delta \mathbf{A}), \quad \bar{\mathbf{B}} = \mathbf{A}^{-1}(\bar{\mathbf{A}} - \mathbf{I})\mathbf{B}$$

Converts continuous dynamics to discrete time steps of size  $\Delta$

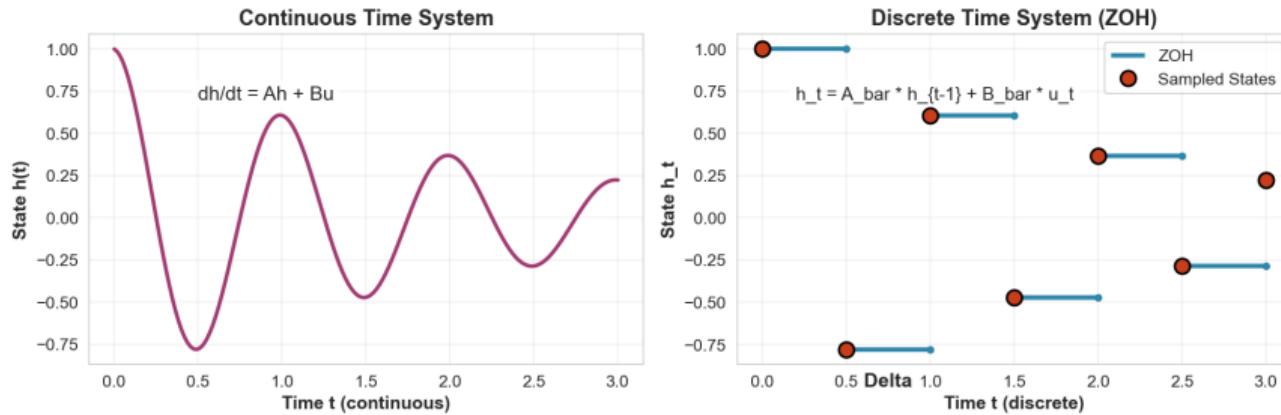
### Discrete Recurrence

$$h_t = \bar{\mathbf{A}}h_{t-1} + \bar{\mathbf{B}}u_t \quad (\text{State update})$$

$$y_t = \mathbf{C}h_t \quad (\text{Output projection})$$

Enables efficient sequential processing of sampled data

## Section Summary: SSM Mathematics (2/2)



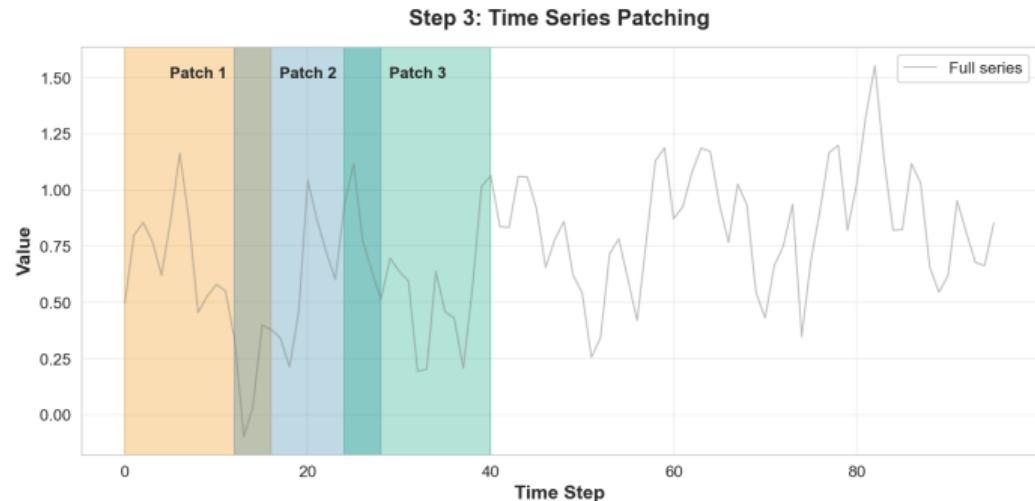
**Key Concept:** Zero-Order Hold (ZOH) maintains input constant between samples

## Step 3: Time Series Patching

The raw time series is typically long and noisy.

We extract **overlapping patches** of fixed length (e.g., 96 steps).

This is analogous to image patching in Vision Transformers.

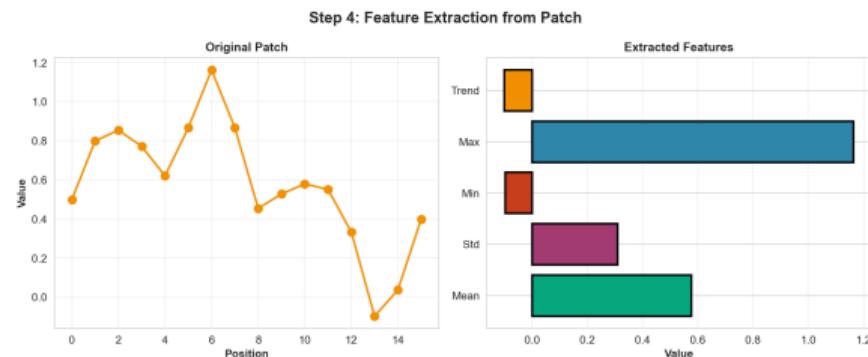


## Step 4: Feature Extraction

From each patch, we can extract:

- **Raw values:** Direct encoding
- **Statistical features:** Mean, variance, etc.
- **Structural features:** Recurrence structure

The Visual Encoder uses the last option.



## Step 5: Visual Transformation (RP)

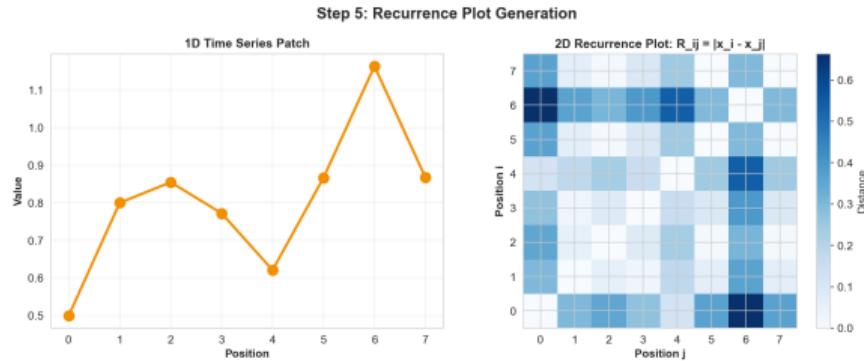
We convert the 1D patch into a 2D Recurrence Plot.

$$R_{i,j} = \|\mathbf{x}_i - \mathbf{x}_j\|$$

**Numerical Example:** Sequence patch (first 2 values):  $\mathbf{x} = [0.5, 0.8]$

$$\mathbf{R} = \begin{pmatrix} |0.5 - 0.5| & |0.5 - 0.8| \\ |0.8 - 0.5| & |0.8 - 0.8| \end{pmatrix}$$

$$\mathbf{R} = \begin{pmatrix} 0 & 0.3 \\ 0.3 & 0 \end{pmatrix}$$



# Algorithm: Selective Scan

Logic verified in: `src/models/mamba_block.py`

---

## Algorithm 1 Selective Scan with HiPPO-Initialized SSM

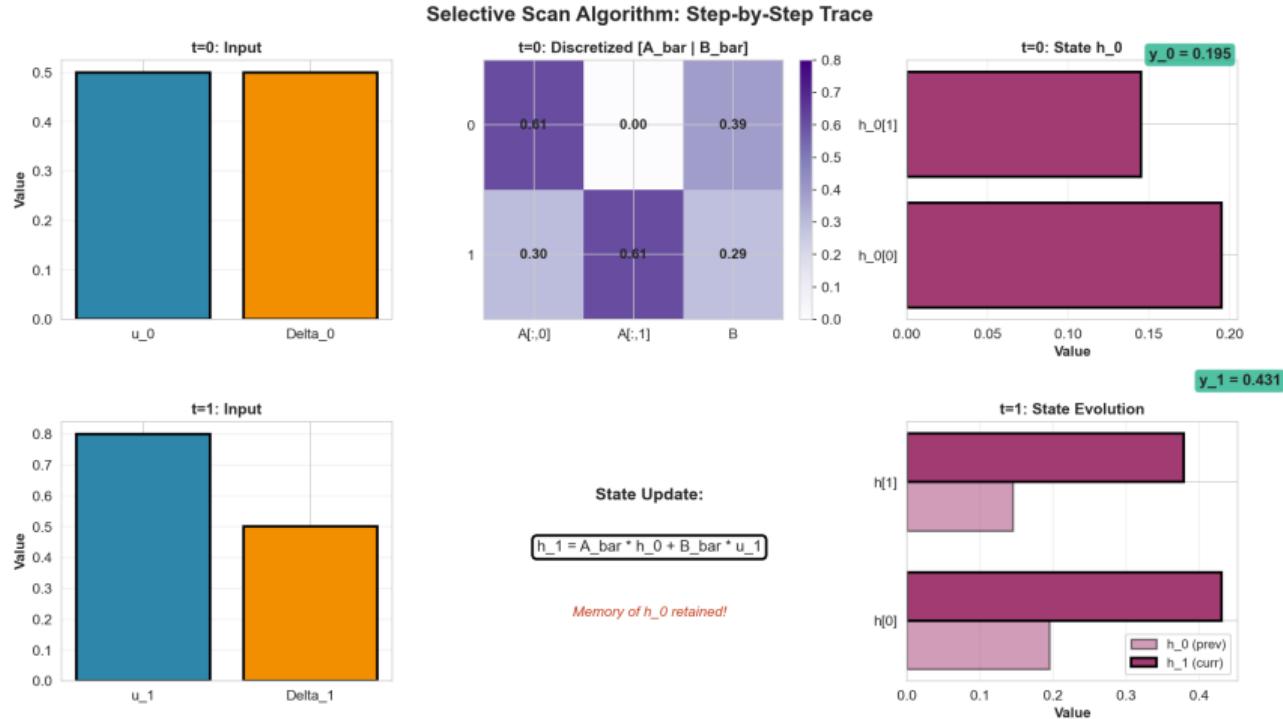
---

Require: Input  $\mathbf{x} \in \mathbb{R}^{B \times T \times D}$ , step sizes  $\boldsymbol{\delta} \in \mathbb{R}^{B \times T \times 1}$

Ensure: Output  $\mathbf{y} \in \mathbb{R}^{B \times T \times D}$

- 1: Precompute  $\mathbf{A}, \mathbf{B}, \mathbf{C}$  (HiPPO)
  - 2: **for**  $t = 0$  **to**  $T - 1$  **do**
  - 3:    $\mathbf{u}_t \leftarrow \mathbf{x}[:, t, :]$ ;  $\Delta_t \leftarrow \boldsymbol{\delta}[:, t, :]$
  - 4:    $\tilde{\mathbf{A}} \leftarrow \Delta_t \cdot \mathbf{A}$  {Scale A by time step}
  - 5:    $\mathbf{A}_d \leftarrow \exp(\tilde{\mathbf{A}})$  {Discretize A}
  - 6:    $\mathbf{h} \leftarrow \mathbf{A}_d \mathbf{h} + \mathbf{B}_d \mathbf{u}_t$  {Update State}
  - 7:    $\mathbf{y}_t \leftarrow \mathbf{C} \mathbf{h}$  {Project to Output}
  - 8: **end for**
-

# Numerical Walkthrough: Selective Scan (1/2)



## Numerical Walkthrough: Selective Scan (2/2)

**Key Insight:** At  $t = 1$ , state  $h_1$  retains memory from both  $u_0$  and  $u_1$ !

- **Input Sequence:**

- $u_0 = 0.5, u_1 = 0.8$  (from our running example)
- Step size:  $\Delta = 0.5$

- **State Evolution:**

- $h_0 = \bar{\mathbf{B}}u_0 \approx [0.195, 0.145]^T$
- $h_1 = \bar{\mathbf{A}}h_0 + \bar{\mathbf{B}}u_1 \approx [0.431, 0.318]^T$
- State accumulates information:  $h_1$  depends on both  $u_0$  and  $u_1$

- **Output Sequence:**

- $y_0 = \mathbf{C}h_0 = 0.195$
- $y_1 = \mathbf{C}h_1 = 0.431$

# Implementation Specifics

Based on `mamba_block.py`:

- **HiPPO Initialization:** Matrix  $\mathbf{A}$  is initialized using the Legendre-S (LegS) measure to handle long-term dependencies.
- **Parallelism vs. Scanning:**
  - The Python implementation performs a **sequential loop** (Line 144 in `mamba_block.py`).
  - Optimized CUDA implementations (not currently used) usually perform a parallel associative scan.
- **Numerical Stability:**
  - $\Delta$  is clamped:  $\Delta \in [10^{-4}, 3.0]$ .
  - `torch.linalg.solve` used instead of explicit inverse for  $\mathbf{A}^{-1}$  to ensure stability.

# Summary: State Space Models for Time Series

## Core Concepts:

- ① **Motivation:** Efficient long-range dependencies
- ② **SSM Math:** Continuous → Discrete via ZOH
- ③ **Selective Scan:** Dynamic state updates
- ④ **Visual Encoding:** Recurrence Plots for 2D patterns

## Implementation:

- HiPPO initialization for  $\mathbf{A}$
- Parallel training + Sequential inference
- Two encoder variants: Standard & Visual
- Numerically stable via `torch.linalg.solve`

**Running Example:** Time series ( $T = 96$ ) with  $x_0 = 0.5, x_1 = 0.8$   
⇒ State memory:  $h_1$  encodes both past and present!

# Questions?